

TAU-MULTILINGUAL: BENCHMARKING VOICE AGENTS ACROSS LANGUAGES

Soham Ray¹ Edgard dos Santos Paiva Neto¹ Ruben D. Valenzuela¹ Karthik Narasimhan³ Keshav Dhandhanian¹ Victor Barres²

¹Sierra ²Mercor ³Princeton University

ABSTRACT

English-only benchmarks expose only a narrow slice of voice-agent behavior. We introduce τ -Multilingual, extending τ -Voice to Spanish, Brazilian Portuguese, Hindi, Korean, and Mandarin with native-speaker review. Across 4,500 full-duplex calls and five voice configurations, Spanish, Portuguese, and Hindi remain within 3.2 task-completion points of English, but Korean and Mandarin fall by 16.8 and 11.3 points. The failure modes also vary: Korean systems miss more responses, Mandarin systems interrupt more often, and both struggle with tools and entities. Moreover, Grok leads in task completion despite having the worst call quality, motivating our separate interaction and experience scores. τ -Multilingual provides a foundation for community-built, culturally grounded voice-agent evaluation.

Index Terms— multilingual voice agents, spoken dialogue evaluation, task-oriented dialogue, full-duplex interaction, user simulation

1. INTRODUCTION

Voice agents must serve a multilingual world: English is spoken by only 19% of the world’s population [1], yet leading end-to-end voice-agent benchmarks evaluate English alone [2, 3, 4]. Evaluating multilingual transfer therefore requires more than checking task completion: it must also capture how agents manage the conversation and whether their speech is linguistically natural and faithful.

We introduce τ -Multilingual, extending τ -Voice [2] through three contributions: *LanguageFactory*, a typed, versioned workflow for native-speaker-authored simulators; a native-speaker-reviewed evaluation that measures interaction and utterance experience separately from task completion; and an evaluation of five voice configurations and two text systems across six languages and three domains.

We [release the codebase](#) so the community can build simulated callers that reflect local linguistic and cultural norms rather than merely translating English personas, alongside language-specific audits that expose English-centric behavior in voice agents.

2. RELATED WORK

τ -Voice brings the database-grounded tasks of τ -bench and τ^2 -Bench [12, 13] to full-duplex speech, but it, EVA-Bench, and FDB-v3 are English-only [2, 3, 4]. M3-DuplexBench covers English and Japanese turn-taking without executed tool tasks [5]; VoiceAgentBench evaluates multilingual spoken tool




Fig. 1. Model-level means with worst-to-best language whiskers; Experience excludes English.

Table 1. Related benchmarks. A check mark denotes direct evaluation and a tilde partial coverage. Spk.: spoken; FD: full duplex; Mult.: multilingual; Int.: interaction.

Work	Spk.	FD	Mult.	Task	Int.	Audit
<i>Voice-agent benchmarks</i>						
τ -Voice [2]	✓	✓	–	✓	✓	–
FDB-v3 [4]	✓	✓	–	✓	✓	–
EVA-Bench [3]	✓	~	–	✓	✓	~
M3-DuplexBench [5]	✓	✓	✓	–	✓	–
<i>Multilingual speech and dialogue</i>						
VoiceAgentBench [6]	✓	–	✓	~	–	–
Speech-MASSIVE [7]	✓	–	✓	–	–	–
X-RiSAWOZ [8]	–	–	✓	–	–	–
GlobalWoZ [9]	–	–	✓	–	–	–
<i>Language-quality evaluation</i>						
MQM [10]	–	–	✓	–	–	✓
MENLO [11]	–	–	✓	–	–	✓
τ-Multilingual	✓	✓	✓	✓	✓	✓

use at the query/workflow level [6]; and Speech-MASSIVE evaluates intent and slot prediction from recorded speech [7]. X-RiSAWOZ and GlobalWoZ provide multilingual text dialogue with human-verified or locale-native content [8, 9]; MQM and MENLO motivate factorized, human-reviewed language evaluation [10, 11]. τ -Multilingual combines these lines in live voice calls with tools, environment execution, and native-speaker-authored fluency audits.

3. METHODOLOGY

3.1. Matched Multilingual Benchmark

τ -Voice supplies the English task frames, policies, tools, binary reward, simulated caller, and 200-ms full-duplex runtime [2]. Each language pack localizes the caller instructions, persona,

voice, identity entities, and evaluation rules while preserving the underlying tasks and scoring. The same specification drives both voice and matched text conditions, enabling comparisons by domain, task, system, and reasoning setting. LanguageFactory stores this specification as a typed, versioned artifact; Table 2 illustrates the Hindi pack. Cross-language differences therefore compare complete localized systems rather than an isolated causal effect of language.

Table 2. LanguageFactory components illustrated with the Hindi pack. Entity localization maps the same source identity to locale-specific backend values without changing the task.

Language component	Hindi (India)
Personas	Rishika uses fast Hinglish; Imran uses patient, Urdu-inflected Hindi.
Speech rules	Spell with “R for Rajdhani”; say <i>at</i> and <i>dot</i> in email addresses.
Task arms and entity localization	Allison Reeves → Neha Gupta; phone: 99400 97181.
Audit factors	Customer address (<i>aap</i> vs. <i>tum/tu</i>) and agreement with <i>aap</i> .
Review and freeze	Pin the Rishika and Imran voices; freeze preset <code>multilingual_v1_hindi</code> .

Native speakers review each pack’s language rules, voices, and smoke calls. Automated checks enforce schema, task parity, and persona assignment before release.

3.2. Evaluation Cohorts

The primary descriptive slice contains 4,500 calls: six languages, three domains, five voice systems, and the same 50 tasks per domain. The voice systems are GPT (OpenAI `gpt-realtime-2`) at minimal and xhigh reasoning, Gemini 3.1 Flash Live Preview at minimal and high, and Grok (xAI Grok Voice Think Fast); the caller is `gpt-5.5 xhigh`. For task-success significance tests only, the four GPT and Gemini configurations include one additional trial. We use the short labels GPT, Gemini, and Grok below.

The matched half-duplex text control contains 1,800 calls over the same six languages, domains, 50-task frames, and localized identities, evaluating GPT-5.5 xhigh and Gemini 3.1 Pro high. Because both the model and interaction regime change, this is an end-to-end control rather than a pure acoustic intervention.

The ablation evaluates the same 30 retail tasks in Hindi and Mandarin with GPT xhigh and Gemini high. Against the localized, romanized baseline, *Entity* substitutes source-English identity values, while *Entity + DB* adds native-script values and matched spell-out, with all else fixed.

3.3. Evaluation Measures

Task success is the conversation-level environment-reward pass rate. We introduce two complementary quality scores. *Interaction* is a call-level composite: 100 minus the equal-weight mean of non-response, inappropriate interruption, selectivity error, a monologue over 45 seconds, and validated tool misuse.

Experience is the share of aligned agent utterances passing both Fluency and Speech fidelity. These composites are compact representations only; their intermediate values have no calibrated interpretation. Fluency flags unnatural phrasing. Speech-fidelity failure requires a severity- ≥ 2 delivery finding or a Mandarin tone-induced meaning flip. Language-specific diagnostics remain separate. English lacks a validated Fluency judge, so it enters Task success and Interaction but not Experience.

Comparison scope. We compare systems within languages and estimate matched English-relative transfer for Task success and Interaction. We do not rank languages: each is a complete localized system, and fluency is not cross-language calibrated.

3.4. Human Validation

Released learned judges enter results only when precision and recall exceed .75, F1 exceeds .80, and Cohen’s $\kappa > .60$. For naturalness and language-specific diagnostics, labels require consensus from two native-language annotators. Validation covers tool use, approximately 60 examples per language for naturalness and speech fidelity, and at least 30 per language-specific diagnostic; failing measures remain archived.

Separately, reviewers audit 30 completed calls per language, balanced across domains and systems. They attribute the first critical error to the agent, simulator, or infrastructure and score eight caller-quality dimensions on a four-point scale. This audit checks whether the simulated caller plausibly supports system comparison; it is not used to train or tune the evaluated agents.

3.5. Statistical Analysis

Task-success inference averages the two trials per task for the four GPT/Gemini configurations, then compares each language with English across 150 matched domain–task clusters using paired permutation tests over five contrasts. We report paired BCa 95% intervals from 10,000 resamples. Grok has one trial and appears only in the descriptive estimates in Figure 2(a).

Interaction and Experience use 100,000 paired cluster-level label-swap permutations. Interaction compares each language with English after jointly swapping the five systems and re-aggregating component-specific failure and opportunity counts. Experience compares all ten system pairs after jointly swapping the five localized calls and re-aggregating aligned utterances. All reported confirmatory *p*-values use Holm correction within each family: five language contrasts for Task success and Interaction, and ten system pairs for Experience. The retail ablations are descriptive.

4. RESULTS

We report task completion, Interaction, and Experience separately; Experience is decomposed into Fluency and Speech fidelity. Figure 2 presents the headline system–language result and decomposes it by model and language. We then examine localization failures and model choice across all three outcomes.

(a) Task completion							(b) Interaction						(c) Experience				
	EN	ES	PT	HI	KO	ZH	EN	ES	PT	HI	KO	ZH	ES	PT	HI	KO	ZH
GPT min.	45	46	49	43	26	30	52	46	46	48	38	42	64	69	58	61	75
GPT xhigh	65	58	69	61	29	44	47	40	43	45	33	39	68	74	62	62	71
Gem min.	41	37	38	39	26	28	49	44	43	48	36	36	58	67	47	66	60
Gem high	51	60	59	60	37	50	45	42	45	42	36	34	61	71	49	57	52
Grok	76	78	79	76	75	69	50	47	47	44	39	32	50	60	40	47	63
Overall	56	56	59	56	39	44	49	44	45	45	36	36	60	68	51	59	64

(d) By model									(e) Best by language						Models
Metric	Level	Unit	GPT min.	GPT xhigh	Gem min.	Gem high	Grok	Avg.	EN	ES	PT	HI	KO	ZH	
Task failure	Conv.	%	60	46	65	47	24	48	24 ⁵	22 ⁵	21 ⁵	24 ⁵	25 ⁵	31 ⁵	1 GPT min.
Non-response	Conv.	%	9	28	47	77	10	34	2 ¹	3 ¹	2 ¹	1 ¹	5 ⁵	3 ¹	2 GPT xhigh
Interruption	Conv.	%	80	90	74	79	90	83	63 ³	66 ³	72 ³	61 ³	82 ⁴	85 ¹	3 Gem min.
Selectivity error	Conv.	%	97	96	84	77	97	90	83 ⁴	82 ⁴	77 ⁴	71 ⁴	75 ⁴	73 ⁴	4 Gem high
Monologue	Conv.	%	18	18	7	11	37	18	1 ⁴	11 ³	11 ³	7 ⁴	3 ³	5 ³	5 Grok
Tool misuse	Conv.	%	70	62	77	53	51	62	36 ⁵	32 ⁵	29 ⁵	43 ⁴	75 ⁴	73 ⁴	
Naturalness failure	Utt.	%	32	30	37	38	44	36	--	30 ²	22 ²	35 ²	30 ³	21 ¹	
Speech-fidelity failure	Utt.	%	7	7	8	9	21	10	3 ²	6 ¹	5 ³	6 ³	6 ²	7 ¹	
Latency	Resp.	secs	1.12	1.42	1.36	1.82	1.20	1.38	1.1 ¹	1.0 ¹	1.1 ¹	1.1 ¹	1.0 ⁵	1.1 ¹	

Fig. 2. Headline outcomes and diagnostics. In (a–c), higher is better; color scales normalize within panel, outlines mark language-best systems, Overall averages models, and Experience excludes English. In (d–e), lower is better; bold marks row minima, and the numbered legend keys the language-best models superscripted in (e). Level gives the conversation, utterance, or response scope.

4.1. Task Completion

Task success is broadly preserved in Spanish, Portuguese, and Hindi but not in Korean and Mandarin (Figure 2(a)). The confirmatory English-relative losses for Korean and Mandarin are 20.7 and 13.3 points (both Holm-adjusted $p < .00005$); the other three differences are not significant. The domain pattern is similarly misaligned: telecom is hardest by task success (44.1%) yet yields the highest Experience (63.2%).

Voice–text gap. Pooled text systems exceed voice task completion in every language by 16.8–41.5 points. The smallest gap is Spanish; the largest is Korean, where text reaches 80.3% versus 38.8% voice. Because model and turn-taking also differ, this is an end-to-end voice-system gap rather than a purely acoustic effect.

4.2. Interaction

Interaction shows the same broad split but a different inferential result: all five English-relative declines are significant ($p \leq .00055$), with substantially larger effects for Korean and Mandarin. Figure 2(e) decomposes the primary result by failure mode and language. The Korean gap concentrates in non-response (46.1% versus 27.8% in English) and tool-use failure (84.9% versus 49.8%). Mandarin instead combines elevated interruption (91.7% versus 78.0%) with tool-use failure (78.8%). Selectivity failure is high in every language (86.3–

93.5%), so it is a broad system weakness rather than the main explanation of the language gap. Spanish and Portuguese show more monologuing than English, but not the Korean/Mandarin non-response and tool-use pattern. Similar aggregate Interaction losses thus arise from different operational failures and imply different remedies.

4.3. Experience

Experience produces a different model ranking: GPT xhigh leads Spanish, Portuguese, and Hindi; Gemini minimal leads Korean; and GPT minimal leads Mandarin (Figure 2(c)).

Fluency and cultural grounding. Native reviewers found recurring literal or interface-like language: Spanish `user id`, Portuguese “nenhum voo ainda foi voado” (no flight has yet been flown), Hindi “explicit yes,” Korean spoken tool names such as `check_sim_status`, and Mandarin `toggle data`. Targeted audits capture different conventions: mixed formal and informal address in Portuguese; gendered self-reference and omitted *-ji* in Hindi; missing subject honorification in Korean; and unsuitable titles and missing sentence-final particles in Mandarin. These diagnostics contextualize Fluency but are not components of Experience. GPT is more natural in Hindi than Gemini and Grok (35.1–40.6% versus 49.7–55.5% failure), yet GPT minimal has the most gender errors (49.3%). In Mandarin, GPT has nearly all modal-particle

failures (40.0–59.1% versus 0–2.0%).

A related, separate diagnostic is gender agreement. Agent gender was inadvertently absent from the prompt despite using female voices, yielding inconsistent self-grounding. Hindi agents default to masculine self-reference (49.5% of calls versus 5.3% with any feminine form). The overall gender-agreement failure rates—27.2% in Hindi and 29.6% in Portuguese—combine self-reference, noun/anaphora agreement, and caller-directed errors. Disclosure changes the mix but does not reliably reduce the total: in 30-call-per-language ablations, failures rise from 30.0% to 40.0% in Hindi and fall from 30.7% to 20.0% in Portuguese.

Speech fidelity. The all-language diagnostic reveals a provider pattern: Grok’s cleanliness is 79.7%, versus 91.4–93.9% for the other systems. Across 71,413 utterances, the most common material failures are punctuation or formatting errors (3,529 utterances), mispronunciations (1,362), and number, date, or currency errors (1,068).

4.4. Localization and Entity Handling

The large Korean and Mandarin task losses concentrate in authentication: exact lookup succeeds in 81–85% of English calls, but only 17–55% of Korean and 37–58% of Mandarin calls across repeated GPT/Gemini systems. Localized caller prompts provide the correct family-first name, yet agents often map the spoken parts into a given-first lookup. In Mandarin, this can break exact lookup even when callers provide the correct romanized spelling and the agent prompt requires all tool calls, arguments, and structured outputs in English; native-character explanations add a separate cross-script mismatch. Telecom’s balanced split shows an exploratory crossover: Korean task completion is 32.0% for name+DOB versus 35.2% for phone authentication, whereas Mandarin is 40.8% versus 33.6%; method remains bundled with task.

Table 3. Retail entity ablations: task completion (%) on 30 tasks.

Lang.	System	Entity		Entity + DB	
		Localized	Source	Romanized	Native
Hindi	GPT xhigh	63.3	63.3	63.3	36.7
	Gemini high	60.0	66.7	60.0	60.0
Mandarin	GPT xhigh	30.0	46.7	30.0	46.7
	Gemini high	43.3	53.3	43.3	60.0

Localization ablations. Source-English entities match or outperform localized ones in every cell of Table 3, improving Mandarin by 16.7 points for GPT and 10.0 for Gemini. Replacing romanized database values with native-script names improves Mandarin by 16.7 points for both systems. Hindi differs: native script leaves Gemini unchanged and lowers GPT by 26.7 points. These 30-task cells are exploratory, and exact authentication does not guarantee downstream task success.

4.5. Model Choice Across Outcomes

Figure 2(d) decomposes the model-level tradeoffs. No configuration dominates; model choice should follow the deployment

objective:

- **Task completion:** Grok reaches 75.6% and leads every language, but has the lowest Experience (51.8%).
- **Quality:** GPT xhigh leads Experience (67.4%) with 54.4% task success; GPT minimal leads Interaction (45.3%) and has the lowest latency (1.12 s).
- **Stability:** Grok varies least in task success (9.4-point range), Gemini high in Interaction (11.5), and GPT xhigh in Experience (12.3); Grok varies most in both quality scores.
- **Latency:** Higher reasoning raises response latency (GPT: 1.12–1.42 s; Gemini: 1.36–1.82 s). GPT/Gemini calls take 11.3–13.8 min, versus 25.6 min for Grok.

4.6. Validation and Robustness

Pooled held-out F1 is 96% for tool use and 88% for both naturalness and speech fidelity; the remaining Interaction components are deterministic. Among measures clearing the frozen gate, naturalness spans Hindi (F1 81.4%, $\kappa = .634$) to Spanish (F1 97.9%, $\kappa = .965$), and speech fidelity spans Korean (F1 82.4%, $\kappa = .754$) to Spanish (F1 96.6%, $\kappa = .955$).

The separate caller audit also supports the evaluation contract. Reviewers attribute only 5/180 first critical errors (2.78%) and 5/104 failed calls (4.81%) to the simulator, and none to infrastructure. The eight caller-quality dimensions average 3.05/4; backchannel naturalness is lowest at 2.86.

Every pair of systems differs significantly on utterance Experience; the weakest contrast is Gemini minimal versus high ($p = .01844$). Interaction effect sizes preserve the substantive distinction between modest Spanish/Portuguese/Hindi degradation and large Korean/Mandarin degradation.

5. LIMITATIONS

We study six languages, three domains, three model families, and one speech stack. Simulated callers may not reproduce human timing or repair. We inject caller text to minimize user-side ASR noise; the simulator likewise consumes agent transcripts, so synthesis errors cannot affect task outcomes. GPT and Gemini have two trials versus Grok’s one; validation sets are finite; and the voice–text and localization ablations change bundled conditions, so their effects remain exploratory.

6. CONCLUSION

τ -Multilingual shows that multilingual voice-agent reliability extends beyond task completion. Spanish, Portuguese, and Hindi largely preserve English task success, whereas Korean and Mandarin lose ground through authentication, entity, and interaction failures. Model rankings shift with the objective, so joint reporting of task success, interaction, and experience exposes failures hidden by English-only or outcome-only benchmarks. We release the benchmark for community-built language packs, audits, and results.

Acknowledgment. We thank Ben Shi, Vijay Iyengar, Ola Zyteck, and Ajeet Grewal for fruitful discussions, and Clay Bavor for his continued support.

7. COMPLIANCE WITH ETHICAL STANDARDS

The benchmark uses no real customer records or recordings. Records are synthetic and non-sensitive; calls use simulated callers and provider voices. Reviewers authored benchmark materials and annotated outputs; they consented and were compensated. Before release, artifacts were screened for incidental personal data or secrets, and flagged items were excluded.

The benchmark evaluates voice-agent systems, not people or languages; its results should not be interpreted as measures of human ability or cultural quality.

This work was supported by Sierra and Mercor. Authors affiliated with Sierra or Mercor are employees of their respective organizations; the authors have no other relevant interests to disclose.

8. REFERENCES

- [1] World Bank, “Digital progress and trends report 2025: Strengthening AI foundations,” Tech. Rep., World Bank, 2025.
- [2] Soham Ray, Keshav Dhandhania, Victor Barres, and Karthik Narasimhan, “ τ -Voice: Benchmarking Full-Duplex Voice Agents on Real-World Domains,” *arXiv preprint arXiv:2603.13686*, 2026.
- [3] Tara Bogavelli, Gabrielle Gauthier Melançon, Katrina Stankiewicz, Oluwanifemi Bamgbose, Fanny Riols, Hoang H. Nguyen, Raghav Mehndiratta, Lindsay Devon Brin, Joseph Marinier, Hari Subramani, Anil Madamala, Sridhar Krishna Nemala, and Srinivas Sunkara, “EVA-Bench: A new end-to-end framework for evaluating voice agents,” *arXiv preprint arXiv:2605.13841*, 2026.
- [4] Guan-Ting Lin, Chen Chen, Zhehuai Chen, and Hung-yi Lee, “Full-duplex-bench-v3: Benchmarking tool use for full-duplex voice agents under real-world disfluency,” *arXiv preprint arXiv:2604.04847*, 2026.
- [5] Ryo Fukuda, Atsushi Ando, Hiroki Kanagawa, Takatomo Kano, Marc Delcroix, Naohiro Tawara, and Yuya Chiba, “M3-DuplexBench: A multi-turn, multilingual, multidomain benchmark for full-duplex spoken dialogue models,” *arXiv preprint arXiv:2607.29125*, 2026.
- [6] Dhruv Jain, Harshit Shukla, Gautam Rajeev, Ashish Kulkarni, Chandra Khatri, and Shubham Agarwal, “VoiceAgentBench: Are voice assistants ready for agentic tasks?,” *arXiv preprint arXiv:2510.07978*, 2025.
- [7] Beomseok Lee, Ioan Calapodescu, Marco Gaido, Matteo Negri, and Laurent Besacier, “Speech-MASSIVE: A multilingual speech dataset for SLU and beyond,” in *Interspeech 2024*, 2024, pp. 817–821.
- [8] Mehrad Moradshahi, Tianhao Shen, Kalika Bali, Monojit Choudhury, Gael de Chalendar, Anmol Goel, Sungkyun Kim, Prashant Kodali, Ponnurangam Kumaraguru, Nasredine Semmar, Sina J. Semnani, Jiwon Seo, Vivek Seshadri, Manish Shrivastava, Michael Sun, Aditya Yadavalli, Chaobin You, Deyi Xiong, and Monica S. Lam, “X-RiSAWOZ: High-quality end-to-end multilingual dialogue datasets and few-shot agents,” in *Findings of the Association for Computational Linguistics: ACL 2023*, 2023, pp. 2773–2794, Association for Computational Linguistics.
- [9] Bosheng Ding, Junjie Hu, Lidong Bing, Mahani Aljunied, Shafiq Joty, Luo Si, and Chunyan Miao, “GlobalWoZ: Globalizing MultiWoZ to develop multilingual task-oriented dialogue systems,” in *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 2022, pp. 1639–1657, Association for Computational Linguistics.
- [10] Markus Freitag, George Foster, David Grangier, Viresh Ratnakar, Qijun Tan, and Wolfgang Macherey, “Experts, errors, and context: A large-scale study of human evaluation for machine translation,” *Transactions of the Association for Computational Linguistics*, vol. 9, pp. 1460–1474, 2021.
- [11] Chenxi Whitehouse, Sebastian Ruder, Tony Zhiyang Lin, Oksana Kurylo, Haruka Takagi, Janice Lam, Nicolò Busetto, and Denise Diaz, “MENLO: From preferences to proficiency—evaluating and modeling native-like quality across 47 languages,” in *International Conference on Learning Representations*, 2026.
- [12] Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik R. Narasimhan, “ τ -bench: A Benchmark for Tool-Agent-User Interaction in Real-World Domains,” in *The Thirteenth International Conference on Learning Representations*, 2025.
- [13] Victor Barres, Honghua Dong, Soham Ray, Xujie Si, and Karthik Narasimhan, “ τ^2 -Bench: Evaluating Conversational Agents in a Dual-Control Environment,” *arXiv preprint arXiv:2506.07982*, 2025.